

distribution_tests

February 5, 2025

1 Using SciPy to Determine Distribution Types

In this tutorial, we'll:

1. Generate random datasets from various distributions.
2. Visualize the datasets using Seaborn to understand their shapes.
3. Use statistical tests from `scipy.stats` to determine if the data fits certain distributions.

Please see <https://docs.scipy.org/doc/scipy/reference/stats.html> for documentation on the SciPy stats module.

You may find this article on statistical distributions a helpful introduction: <https://datasciencedojo.com/blog/types-of-statistical-distributions-in-ml/>

```
[1]: # Import NumPy for numerical operations
import numpy as np

# Import statistical functions and distributions from SciPy
from scipy import stats

# Import Seaborn for data visualization
import seaborn as sns

# Set Seaborn style for better aesthetics
sns.set_theme(style='whitegrid')
```

1.1 Generate Random Datasets

We'll create datasets from the following distributions: - Normal Distribution - Exponential Distribution - Poisson Distribution - Uniform Distribution


```
[2]: # First we set the random seed for reproducibility. This will ensure that we
      ↪ get the
      # same randomly generated datasets each time we run the code.
      np.random.seed(42)
```

1.1.1 Normal Distribution

```
[3]: # Generate 1000 random numbers from a normal distribution with mean=0 and std=1
data_normal = np.random.normal(loc=0, scale=1, size=1000)
```

```
[4]: sns.histplot(data_normal)
```

```
[4]: <Axes: ylabel='Count'>
```




1.1.2 Exponential Distribution

```
[5]: # Generate 1000 random numbers from an exponential distribution with scale=1
data_exponential = np.random.exponential(scale=1, size=1000)
```

```
[6]: sns.histplot(data_exponential)
```

```
[6]: <Axes: ylabel='Count'>
```




1.1.3 Poisson Distribution

```
[7]: # Generate 1000 random numbers from a Poisson distribution with lambda=3
data_poisson = np.random.poisson(lam=3, size=1000)
```

```
[8]: sns.histplot(data_poisson)
```

```
[8]: <Axes: ylabel='Count'>
```




1.1.4 Uniform Distribution

```
[9]: # Generate 1000 random numbers from a uniform distribution between -2 and 2
data_uniform = np.random.uniform(low=-2, high=2, size=1000)
```

```
[10]: sns.histplot(data_uniform)
```

```
[10]: <Axes: ylabel='Count'>
```



1.2 Statistical Tests to Determine Distribution Type

We will use the following statistical tests:

- Normality Tests: To check if data follows a normal distribution.
 - Shapiro-Wilk Test
 - Kolmogorov-Smirnov Test
- Exponential Distribution Test:
 - Kolmogorov-Smirnov Test
- Poisson Distribution Test:
 - Chi-Squared Goodness of Fit Test
- Uniform Distribution Test:
 - Kolmogorov-Smirnov Test

1.2.1 Checking for Normality

Shapiro-Wilk Test

```
[11]: def shapiro_wilk_test(data):
      """
      Performs the Shapiro-Wilk test for normality. This function is really just
      a wrapper for stats.shapiro().

      Parameters:
```

```
- data: array-like, the dataset to test
```

```
Returns:
```

```
- W statistic and p-value
```

```
"""
```

```
w_stat, p = stats.shapiro(data)
```

```
return w_stat, p
```

```
[12]: # Test normality of data_normal
stat, p = shapiro_wilk_test(data_normal)

print('Shapiro-Wilk Test for Normal Distribution Data:')
print(f'W Statistic: {stat:.4f}, p-value: {p:.4f}')

if p > 0.05:
    print('Fail to reject the null hypothesis: Data is normally distributed.')
else:
    print('Reject the null hypothesis: Data is not normally distributed.')
```

Shapiro-Wilk Test for Normal Distribution Data:

W Statistic: 0.9986, p-value: 0.6273

Fail to reject the null hypothesis: Data is normally distributed.

Interpretation: High p-value (> 0.05): Fail to reject the null hypothesis (data follows the distribution). Low p-value (< 0.05): Reject the null hypothesis (data does not follow the distribution).

Kolmogorov-Smirnov Test for Normality

```
[13]: def ks_test_normal(data):
    """
    Performs the Kolmogorov-Smirnov test for normality after standardizing the
    ↪ data.
    Standardizing transforms the data to have a mean of 0 and a standard
    ↪ deviation of 1
    This ensures that the data is on the same scale as the normal distribution,
    ↪ which
    the test uses as a reference. This makes the test more accurate when used
    ↪ with
    datasets created with different scales or units.

    Parameters:
    - data: array-like, the dataset to test

    Returns:
    - D statistic and p-value
    """
    standardized_data = (data - np.mean(data)) / np.std(data)
```

```
d_stat, p = stats.kstest(standardized_data, 'norm')
return d_stat, p
```

```
[14]: # Test normality of data_normal
stat, p = ks_test_normal(data_normal)

print('Kolmogorov-Smirnov Test for Normal Distribution Data:')
print(f'D Statistic: {stat:.4f}, p-value: {p:.4f}')

if p > 0.05:
    print('Fail to reject the null hypothesis: Data is normally distributed.')
else:
    print('Reject the null hypothesis: Data is not normally distributed.')
```

Kolmogorov-Smirnov Test for Normal Distribution Data:
D Statistic: 0.0215, p-value: 0.7370
Fail to reject the null hypothesis: Data is normally distributed.

1.2.2 Testing for Exponential Distribution

```
[15]: def ks_test_exponential(data):
    """
    Performs the Kolmogorov-Smirnov test for exponential distribution.

    Parameters:
    - data: array-like, the dataset to test

    Returns:
    - D statistic and p-value
    """
    # Shift data to ensure all values are >= 0. The exponential distribution is
    →defined
    # for values >= 0. If the data contains negative values, the test could
    →produce inaccurate
    # results or fail.
    shifted_data = data - np.min(data)

    # Calculate loc and scale, allowing the test to tailor the exponential
    →distribution to the
    # data. This makes the comparison more meaningful and accurate.
    loc, scale = stats.expon.fit(shifted_data)

    d_stat, p = stats.kstest(shifted_data, 'expon', args=(loc, scale))
    return d_stat, p
```

```
[16]: # Test if data_exponential follows an exponential distribution
stat, p = ks_test_exponential(data_exponential)
```

```

print('Kolmogorov-Smirnov Test for Exponential Distribution Data:')
print(f'D Statistic: {stat:.4f}, p-value: {p:.4f}')

if p > 0.05:
    print('Fail to reject the null hypothesis: Data follows an exponential_
    ↪distribution.')
else:
    print('Reject the null hypothesis: Data does not follow an exponential_
    ↪distribution.')

```

Kolmogorov-Smirnov Test for Exponential Distribution Data:

D Statistic: 0.0203, p-value: 0.7965

Fail to reject the null hypothesis: Data follows an exponential distribution.

1.2.3 Testing for Poisson Distribution

```

[17]: def chi_square_test_poisson(data):
        """
        Performs the Chi-Squared test for Poisson distribution.

        Parameters:
        - data: array-like, the dataset to test

        Returns:
        - Chi-squared statistic and p-value
        """
        # Get the maximum observed value to define bins. Bins should extend up to_
        ↪the
        # maximum value to ensure that the entire range of the data is represented,
        # which is essential for accurate frequency calculation.
        max_value = np.max(data)

        # Create bins for each integer value up to max_value. This step defines the_
        ↪bins
        # for the histogram to count observed frequencies. The - 0.5 shift centers_
        ↪the bins
        # around integer values, ensuring that each integer is correctly counted in_
        ↪its
        # respective bin. This alignment helps in accurately calculating the_
        ↪observed
        # frequencies for a Poisson distribution.
        bins = np.arange(0, max_value + 2) - 0.5

        # Calculate observed frequencies using histogram. The histogram function is_
        ↪used to

```



```

# calculate the number of data points that fall into each bin, effectively
↳computing
# the observed frequency distribution of the data. This step is necessary
↳for comparing
# observed counts with expected counts in the chi-squared test.
observed_freq, _ = np.histogram(data, bins=bins)

# Estimate the lambda parameter (mean of the data). The mean of the data
↳serves as an
# estimate for the parameter lambda ( ) in a Poisson distribution. The
↳Poisson
# distribution is defined by its mean, so estimating lambda from the data
↳allows for
# the generation of an expected distribution that matches the observed
↳data's average
# rate of occurrence.
lambda_estimate = np.mean(data)

# Calculate expected frequencies using the Poisson PMF. This step computes
↳the expected
# frequency for each integer value using the Poisson probability mass
↳function (PMF).
# By multiplying the PMF values by the total number of data points, the
↳function scales
# the probabilities to match the size of the dataset, resulting in expected
↳frequencies
# for comparison against the observed data.
expected_counts = stats.poisson.pmf(
    np.arange(0, max_value + 1), lambda_estimate) * len(data)

# Adjust expected frequencies to ensure the sums match. This adjustment
↳ensures that the
# total sum of the expected frequencies matches the total number of
↳observed data points.
# Due to potential rounding errors or differences in distribution tails,
↳this
# normalization helps align the expected and observed total counts,
↳improving the accuracy
# of the chi-squared test.
expected_freq = expected_counts * \
    (observed_freq.sum() / expected_counts.sum())

# Perform the Chi-Squared test. The chi-squared test is performed by
↳comparing the
# observed and expected frequency distributions. The test returns a
↳chi-squared statistic,

```

```

    # which quantifies the difference between the distributions, and a p-value,
    ↪which
    # indicates the significance of that difference. This final step allows for
    ↪hypothesis
    # testing to determine if the observed data follows a Poisson distribution.
    chi_squared_stat, p = stats.chisquare(f_obs=observed_freq,
    ↪f_exp=expected_freq)
    return chi_squared_stat, p

```

```

[18]: # Test if data_poisson follows a Poisson distribution
stat, p = chi_square_test_poisson(data_poisson)

print('Chi-Squared Test for Poisson Distribution Data:')
print(f'Chi-Squared Statistic: {stat:.4f}, p-value: {p:.4f}')

if p > 0.05:
    print('Fail to reject the null hypothesis: Data follows a Poisson
    ↪distribution.')
else:
    print('Reject the null hypothesis: Data does not follow a Poisson
    ↪distribution.')

```

Chi-Squared Test for Poisson Distribution Data:

Chi-Squared Statistic: 8.6727, p-value: 0.4680

Fail to reject the null hypothesis: Data follows a Poisson distribution.

1.2.4 Testing for Uniform Distribution

```

[19]: def ks_test_uniform(data):
    """
    Performs the Kolmogorov-Smirnov test for uniform distribution.

    Parameters:
    - data: array-like, the dataset to test

    Returns:
    - D statistic and p-value
    """
    # This step standardizes the data so that it fits within the range [0, 1].
    # The transformation rescales the minimum value of data to 0 and the maximum
    # value to 1, with all intermediate values proportionally adjusted within
    # this range.
    # The Kolmogorov-Smirnov (KS) test checks the similarity between the
    ↪empirical
    # distribution of the data and a specified distribution (in this case, a
    # uniform distribution). A uniform distribution over the interval [0, 1]
    # assumes that all data points are equally likely across this range. By

```

```

# scaling the data to [0, 1], the comparison becomes valid and meaningful.
scaled_data = (data - np.min(data)) / (np.max(data) - np.min(data))

stat, p = stats.kstest(scaled_data, 'uniform')
return stat, p

```

```

[20]: # Test if data_uniform follows a uniform distribution
stat, p = ks_test_uniform(data_uniform)

print('Kolmogorov-Smirnov Test for Uniform Distribution Data:')
print(f'D Statistic: {stat:.4f}, p-value: {p:.4f}')

if p > 0.05:
    print('Fail to reject the null hypothesis: Data follows a uniform_
↪distribution.')
else:
    print('Reject the null hypothesis: Data does not follow a uniform_
↪distribution.')

```

Kolmogorov-Smirnov Test for Uniform Distribution Data:

D Statistic: 0.0263, p-value: 0.4838

Fail to reject the null hypothesis: Data follows a uniform distribution.

1.3 What about stat?

You may have noticed that each of our test functions returns both a stat and a p-value. We only used the p-value in the examples above. The stat value supplies additional context. Below we show an example that uses the W statistic returned by the Shapiro-Wilk test for normality as well as the p-value. To show how this varies, we will use normally distributed data that has been skewed.

```

[21]: # Function to generate skewed data
def generate_skewed_data(size, skewness):
    """
    Generates data with a specified skewness.

    Parameters:
    - size: int, number of data points
    - skewness: float, skewness parameter

    Returns:
    - data: array-like, generated dataset
    """
    # Generate data from a normal distribution
    data = np.random.normal(loc=0, scale=1, size=size)

    # Apply exponential transformation to introduce skewness
    data_skewed = np.sign(data) * np.abs(data) ** (1 + skewness)
    return data_skewed


```

Let's visualise the effects of skew as defined in the function above. We'll see that, although they are all based on a normal distribution, they diverge further and further as we increase skewness.

```
[22]: data_normal = generate_skewed_data(1000, skewness=0.0)
      data_moderate_skew = generate_skewed_data(1000, skewness=0.5)
      data_high_skew = generate_skewed_data(1000, skewness=1.0)
```


```
[23]: # Normally distributed data with skewness=0.0.
      sns.histplot(data_normal)
```

```
[23]: <Axes: ylabel='Count'>
```




```
[24]: # Normally distributed data with skewness=0.5.
      sns.histplot(data_moderate_skew)
```

```
[24]: <Axes: ylabel='Count'>
```



```
[25]: # Normally distributed data with skewness=1.0.  
sns.histplot(data_high_skew)
```

```
[25]: <Axes: ylabel='Count'>
```



```
[26]: # Test normality of the unskewed data
stat, p = shapiro_wilk_test(data_normal)

stat, p
```

[26]: (0.998820161474999, 0.767897403236018)

Interpretation of W Statistic: > 0.95: data closely resembles a normal distribution. Interpretation of p-value: > 0.05: fail to reject the null hypothesis (data follows the normal distribution).

```
[27]: # Test normality of the moderately skewed data
stat, p = shapiro_wilk_test(data_moderate_skew)

print(f'{stat:.2f}, {p:.20f}')
stat, p
```

0.96, 0.000000000000000071183

[27]: (0.9601665499898253, 7.118311417809235e-16)

Interpretation of W Statistic: 0.90 to 0.95: data somewhat resembles a normal distribution. Interpretation of p-value: < 0.05: reject the null hypothesis (data is not normally distributed).

0.85, 0.000000000000000000000000000000002008151

Interpretation of W Statistic: < 0.90 : data does not resemble a normal distribution. Interpretation of p-value: < 0.05 : reject the null hypothesis (data is not normally distributed).

15